

Computational Photography

Instructor: Sanjeev J. Koppal

MWF

1145am-1235pm

BEN 328

Acknowledgements

Some slides from
Narasimhan (Carnegie Mellon),
Zickler (Harvard),
and
Efros (Berkeley)

Assignment updates

What is an image?

We can think of an **image** as a function, f , from \mathbb{R}^2 to \mathbb{R} :

- $f(x, y)$ gives the **intensity** at position (x, y)
- Realistically, we expect the image only to be defined over a rectangle, with a finite range:
 - $f: [a,b] \times [c,d] \rightarrow [0,1]$



A color image is just three functions pasted together.
We can write this as a “vector-valued” function:

$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

Participation

Can you think of something that
would be in a mathematical model
for $f(x,y)$?

Images as functions



What is a digital image?

We usually operate on **digital (discrete)** images:

- **Sample** the 2D space on a regular grid
- **Quantize** each sample (round to nearest integer)

If our samples are \otimes apart, we can write this as:

$$f[i, j] = \text{Quantize}\{ f(i \otimes, j \otimes) \}$$

The image can now be represented as a matrix of integer values

$i \downarrow$ $j \rightarrow$

| | | | | | | | |
|-----|-----|-----|-----|-----|-----|----|-----|
| 62 | 79 | 23 | 119 | 120 | 105 | 4 | 0 |
| 10 | 10 | 9 | 62 | 12 | 78 | 34 | 0 |
| 10 | 58 | 197 | 46 | 46 | 0 | 0 | 48 |
| 176 | 135 | 5 | 188 | 191 | 68 | 0 | 49 |
| 2 | 1 | 1 | 29 | 26 | 37 | 0 | 77 |
| 0 | 89 | 144 | 147 | 187 | 102 | 62 | 208 |
| 255 | 252 | 0 | 166 | 123 | 62 | 0 | 31 |
| 166 | 63 | 127 | 17 | 1 | 0 | 99 | 30 |

Lesson 1

Images can be treated as functions.
Since they are discrete, these discrete functions are represented as matrices.

Image Processing

An **image processing** operation typically defines a new image g in terms of an existing image f .

We can transform either the range of f .

$$g(x, y) = t(f(x, y))$$

Or the domain of f :

$$g(x, y) = f(t_x(x, y), t_y(x, y))$$

What kinds of operations can each perform?

Image Processing

image filtering: change **range** of image

$$g(x) = h(f(x))$$




image warping: change **domain** of image




Image Processing

image filtering: change **range** of image

$$g(x) = h(f(x))$$





image warping: change **domain** of image



Point Processing

The simplest kind of range transformations are these independent of position x,y :

$$g = t(f)$$

This is called point processing.


What can they do?

What's the form of t ?

Important: every pixel for himself – spatial information completely lost!

Basic Point Processing

FIGURE 3.3 Some basic gray-level transformation functions used for image enhancement.



Participation

Could there be other curves?

If so, what kinds?

Point processing

Types of point processing

Negative




Log

a b

FIGURE 3.5

(a) Fourier spectrum.

(b) Result of applying the log transformation given in Eq. (3.2-2) with $c = 1$.



Power-law transformations




FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases).

$$s = cr^\gamma$$

Image Enhancement

| | |
|---|---|
| a | b |
| c | d |

FIGURE 3.9

(a) Aerial image.
(b)–(d) Results of
applying the
transformation in
Eq. (3.2-3) with
 $c = 1$ and
 $\gamma = 3.0, 4.0,$ and
 5.0 , respectively.
(Original image
for this example
courtesy of
NASA.)




Example: Gamma Correction

a b
c d

FIGURE 3.7



(a) Linear-wedge gray-scale image.
(b) Response of monitor to linear wedge.
(c) Gamma-corrected wedge.
(d) Output of monitor.



$$S = r^\gamma$$

e.g. $0.25 = 0.5^{2.0}$

Contrast Stretching

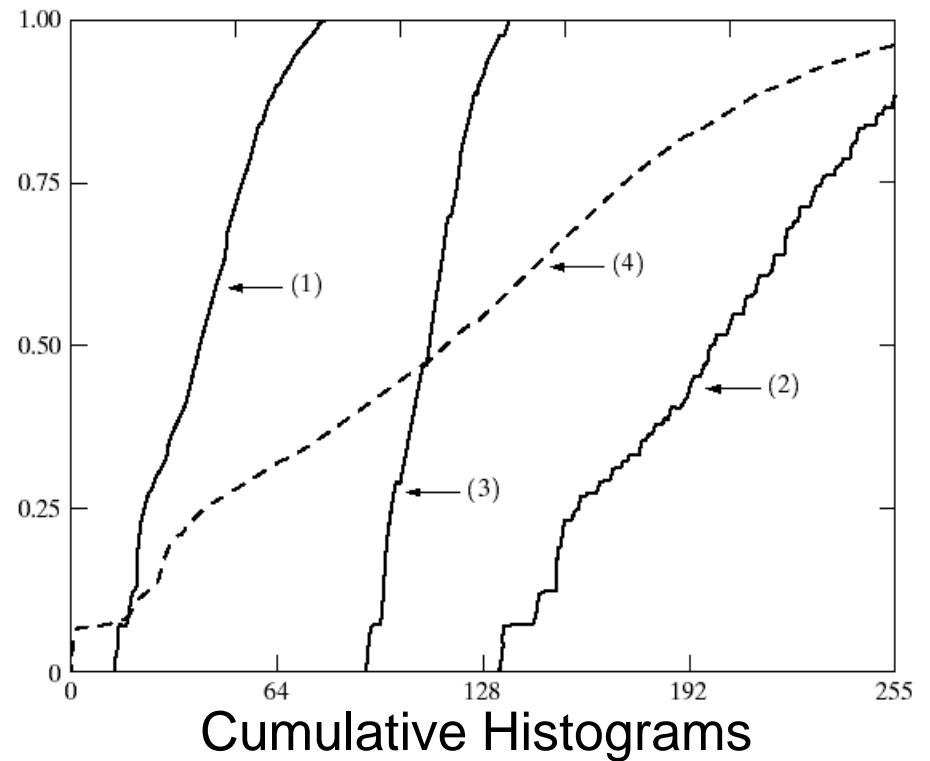



a b
c d

FIGURE 3.10

Contrast stretching.
(a) Form of transformation function. (b) A low-contrast image. (c) Result of contrast stretching. (d) Result of thresholding. (Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)

Image Histograms




$$s = T(r)$$

a b

FIGURE 3.15 Four basic image types: dark, light, low contrast, high contrast, and their corresponding histograms. (Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)

Histogram Equalization



a b c

FIGURE 3.17 (a) Images from Fig. 3.15. (b) Results of histogram equalization. (c) Corresponding histograms.

Neighborhood Processing (filtering)

Q: What happens if I reshuffle all pixels within the image?




Lesson 2

Always understand if the computation being applied to the image is per pixel, patch-based or global.

Neighborhood Processing (filtering)

- Q: What happens if I reshuffle all pixels within the image?



- A: It's histogram won't change. No point processing will be affected...
- Need spatial information to capture this...

Image Filtering

Filtering noise

How can we “smooth” away noise in an image?

| | | | | | | | | | |
|---|---|---|-----|-----|-----|-----|-----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 100 | 130 | 110 | 120 | 110 | 0 | 0 |
| 0 | 0 | 0 | 110 | 90 | 100 | 90 | 100 | 0 | 0 |
| 0 | 0 | 0 | 130 | 100 | 90 | 130 | 110 | 0 | 0 |
| 0 | 0 | 0 | 120 | 100 | 130 | 110 | 120 | 0 | 0 |
| 0 | 0 | 0 | 90 | 110 | 80 | 120 | 100 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Mean filtering

| | | | | | | | | | |
|-----------|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| $F[x, y]$ | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | |
|-----------|--|--|--|--|--|--|--|--|--|
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| $G[x, y]$ | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Mean filtering

| | | | | | | | | | |
|-----------|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| $F[x, y]$ | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | | | | | | | | |
|-----------|----|----|----|----|----|----|----|----|--|
| | | | | | | | | | |
| | 0 | 10 | 20 | 30 | 30 | 30 | 20 | 10 | |
| | 0 | 20 | 40 | 60 | 60 | 60 | 40 | 20 | |
| | 0 | 30 | 60 | 90 | 90 | 90 | 60 | 30 | |
| $G[x, y]$ | 30 | 50 | 80 | 80 | 90 | 60 | 30 | | |
| | 30 | 50 | 80 | 80 | 90 | 60 | 30 | | |
| | 0 | 20 | 30 | 50 | 50 | 60 | 40 | 20 | |
| | 10 | 20 | 30 | 30 | 30 | 30 | 20 | 10 | |
| | 10 | 10 | 10 | 0 | 0 | 0 | 0 | 0 | |
| | | | | | | | | | |

Cross-correlation filtering

Let's write this down as an equation. Assume the averaging window is $(2k+1) \times (2k+1)$:

$$G[i, j] = \frac{1}{(2k+1)^2} \sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]$$

We can generalize this idea by allowing different weights for different neighboring pixels:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v]$$

This is called a **cross-correlation** operation and written:

$$G = H \otimes F$$

H is called the “filter,” “kernel,” or “mask.”

The above allows negative filter indices. When you implement need to use: $H[u+k, v+k]$ instead of $H[u, v]$


Mean kernel

What's the kernel for a 3x3 mean filter?

| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

$F[x, y]$

$H[u, v]$



When can taking an un weighted mean be bad idea?

Gaussian filtering

A Gaussian kernel gives less weight to pixels further from the center of the window


| | | | | | | | | | |
|---|---|----|----|----|----|----|----|---|----|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 1 |
| 0 | 0 | 0 | 90 | 0 | 90 | 90 | 90 | 0 | 16 |
| 0 | 0 | 0 | 90 | 90 | 90 | 90 | 90 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 90 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

| | | |
|---|---|---|
| 1 | 2 | 1 |
| 2 | 4 | 2 |
| 1 | 2 | 1 |


$H[u, v]$

What happens if you increase σ ? $F[x, y]$

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$



Mean vs. Gaussian filtering



Fourier transforms and ALL kinds of 2D signal processing is possible on images.

Fourier transforms and ALL kinds of 2D signal processing is possible on images.

For now we will look at how optics can provide some of these image processing functions.

Lenses


Participation

Straightaway, there is one kind of image processing that we can with lenses.

What is it?




Mean vs. Gaussian filtering




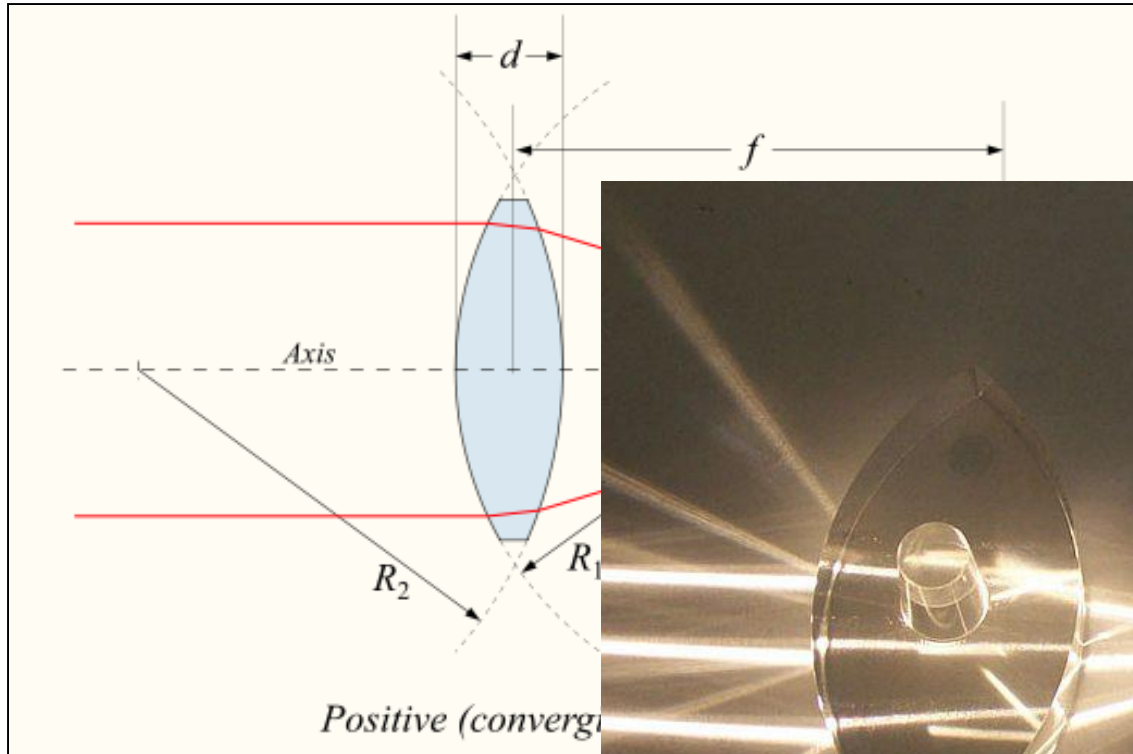
Lesson 3

Optics and code can both perform computations on images.



Types of Lenses



Bi-convex Lens



Bi-concave Lens



Lensmaker's Equation

$$\frac{1}{f} = (n - 1) \left[\frac{1}{R_1} - \frac{1}{R_2} + \frac{(n - 1)d}{n R_1 R_2} \right],$$

f is the focal length of the lens,

n is the refractive index of the lens material,

R_1 is the radius of curvature of the lens surface closest to the light source,


R_2 is the radius of curvature of the lens surface farthest from the light source,

d is the thickness of the lens (the distance along the lens axis between the two surface vertices).

Participation

What about plano-convex lenses?

Optics of a Two Lens System





- Rule : Image formed by first lens is the object for the second lens.
- Main Rays : Ray passing through focus emerges parallel to optical axis.
Ray through optical center passes un-deviated.

- Magnification:
$$m = \frac{i_2}{o_2} \frac{i_1}{o_1}$$

Exercises: What is the combined focal length of the system?
What is the combined focal length if $d = 0$?

Cylindrical Lenses

Circular And Rectangular Cylinder Lenses




Participation

What kind of processing can we do
with this type of lens?


Image Processing

image warping: change ***domain*** of image



Suppose we wanted to scale the ~~image?~~

image warping: change ***domain*** of image



Focusing a Laser beam to a point






Image Processing

image warping: change ***domain*** of image




Changing Diameter of Collimated Beam




Telecentric Lenses

Object-side and Image-side telecentricity:



- Sizes of object and image do not change as they are translated.
- However, focus does change as in any lens.

Eliminating Perspective Distortion




Regular Lens




Telecentric Lens

Illusions with Telecentric Lenses




Relay Lenses (1:1 imaging)

MVO® Relay Lenses- F/4 Design




MVO® Relay Lenses- F/8 Design



Wide angle Lenses




Circular Fisheye



Full Frame Rectangular Fisheye

Vignetting



More light passes through lens L3 for scene point A than scene point B

Results in spatially non-uniform brightness (in the periphery of the image)

Lens Vignetting




- Usually brighter at the center and darker at the periphery.

Vignetting



photo by Robert Johnes



Chromatic Aberrations



Reading: <http://www.dpreview.com>

Lens Glare



- Stray interreflections of light within the optical lens system.
- Happens when very bright sources are present in the scene.

Geometric Lens Distortions





Photo by Helmut Dersch

Both due to lens imperfection
Rectify with geometric camera calibration



Filters

Neutral Density Filters





- Spectrally flat from 400-700nm
- Homogeneous Glass: Blocks by absorption or by reflection
- Light/Exposure Control for Imaging
- Transmittance = $10^{-(\text{optical density})} \times 100$

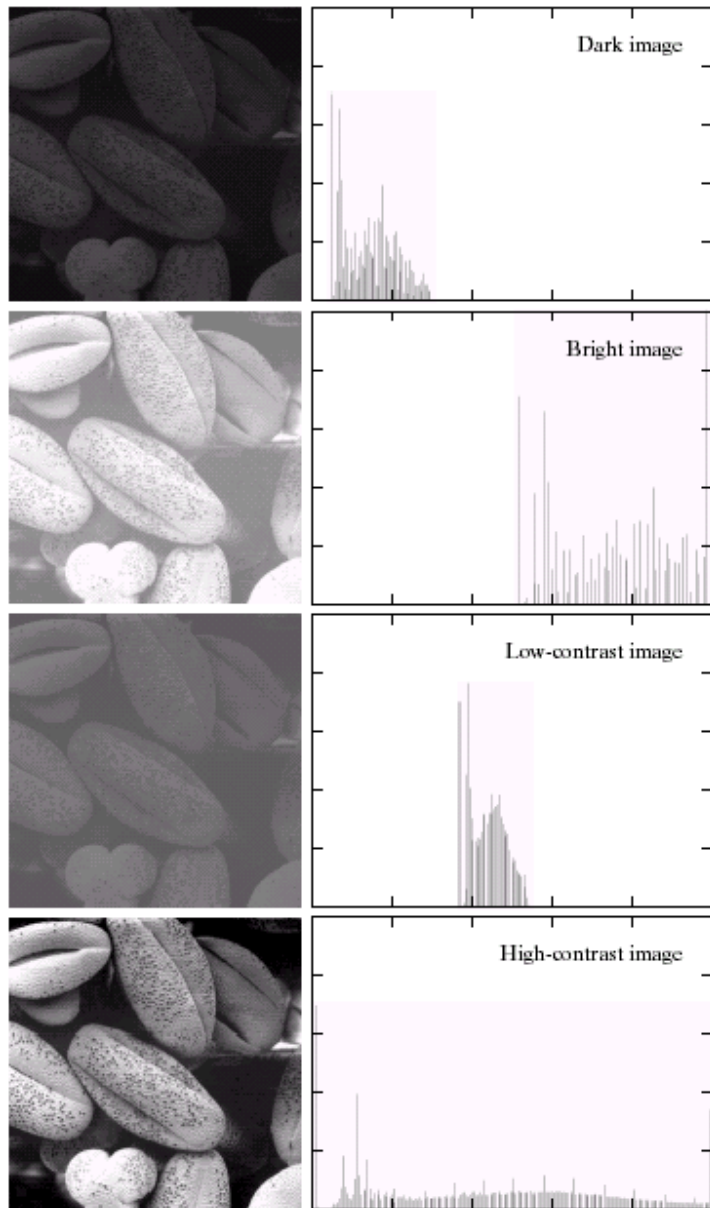
Variable Neutral Density Filters



Variable Neutral Density Filters



Can we get these kinds of effects?






a b

FIGURE 3.15 Four basic image types: dark, light, low contrast, high contrast, and their corresponding histograms. (Original image courtesy of Dr. Roger Heady, Research School of Biological Sciences, Australian National University, Canberra, Australia.)




Show papers

Color Filters



Filter Book





Linear Variable Interference Filters



Mirrors


Cold Mirrors

- Reflect VIS
- Transmit IR




Hot Mirrors



- Reflect IR
- Transmit VIS



Angle of Incidence



Angle of Incidence




Mirrors and Color Filters


Show papers

Prisms can rotate images

Image Reflection



Right Angle Prism



Penta Prism






Image Rotation



Dove Prism



Uncoated



Coated

What can we *not* do with optics (at least non-exotic optics):

- negative numbers**
- histograms (counting and taking max/mins)**
- what else?**

Thank you

Lesson summary:

1. Images are discrete functions that are represented by matrices.
2. Always understand if the computation applied on an image is per pixel, patch-based or global.
3. Optics and code can both perform functions on images